ArchaeoGLOBE Analysis

Nick Gauthier and Ben Marwick

Last knit on: 06 February, 2019

Analysis code for the ArchaeoGlobe database. Here we use Generalized Additive Models (GAMs), a flexible form of nonlinear regression model capable of fitting smooth, time-varying trends to the ordered categorical ArchaeoGLOBE response data.

We model ordered categorical data using a latent variable following a logistic distribution. The model identifies a series of cut points, which correspond the the probabilities of the latent variable falling within each of our categories.

We fit two sets of trends. One trend is fitted to all the data simultaneously, representing the global trend across all archaeological regions. Then we fit region-level trends, which represent the deviation of each region from the global trend. By penalizing the "wiggliness" of the trend lines, we allow regional trends that don't significantly deviate from the global trend to be penalized to 0, effectively reducing that particular region to the global trend. This is a form of partial pooling, allowing the model to share information between groups and in so doing make the results less sensitive to regions with exceptionally low response rates.

After fitting the model, we can extract the region-specific trends, use a k-means clustering algorithm to group together regions with similar trends, and map the results. We repeat this analysis for both self-reported expertise and perceived data quality.

Setup

Import packages needed for analysis. We'll use packages from the `tidyverse`, such as `readr`, `dplyr`, and `ggplot2` for data import, processing, and plotting. We'll also use `mgcv` for fitting nonlinear trends to the data. We'll use the `sf` package to help us plot shapefiles in a tidy context. The `dataverse` and `osfr` packages allows us to pull the raw survey data and precomputed analysis files from their online repositories. Finally, we'll use `patchwork` (installed from GitHub) to combine multiple ggplots in the same image.

```r
library(tidyverse)
library(mgcv)
library(sf)
library(ggplot2)
library(dataverse)

#install patchwork  and osfr from github if needed
#devtools::install_github('thomasp85/patchwork')
library(patchwork)
# devtools::install_github('centerforopenscience/osfr')
library(osfr)
```

Data import

Read in the latest version of the ArchaeoGLOBE database and the consensus assessment from the Dataverse repository.

```r
Sys.setenv('DATAVERSE_SERVER' = 'dataverse.harvard.edu')
```

```r
# get data frame of files on dataverse
ArchaeoGLOBE_Data_DOI <- 'doi:10.7910/DVN/CNCANQ'
ArchaeoGLOBE_Data_df <- get_dataset(ArchaeoGLOBE_Data_DOI)

# Only download the file we need here
ArchaeoGLOBE_Data_df_files <- ArchaeoGLOBE_Data_df$files[
  grepl('ARCHAEOGLOBE_PUBLIC_DATA|ARCHAEOGLOBE_CONSENSUS_ASSESSMENT',
        ArchaeoGLOBE_Data_df$files$filename), ]

# read into local dir
walk(ArchaeoGLOBE_Data_df_files$label,
     ~get_file(.x, ArchaeoGLOBE_Data_DOI) %>%
       writeBin(paste0('data/raw-data/', .x)))

# read into the current environment
archaeoglobe <- read_csv('data/raw-data/ARCHAEOGLOBE_PUBLIC_DATA.tab')
consensus <- read_csv('data/raw-data/ARCHAEOGLOBE_CONSENSUS_ASSESSMENT.tab')
```

Repeat for the archaeological regions shapefile.

```r
# repeat for shapefile
ArchaeoGLOBE_Regions_DOI <- 'doi:10.7910/DVN/CQWUBI'

# get data frame of files on DV
ArchaeoGLOBE_Regions_df <- get_dataset(ArchaeoGLOBE_Regions_DOI)

# just download the shapefile we want
ArchaeoGLOBE_Regions_df_files <- ArchaeoGLOBE_Regions_df$files[
  ArchaeoGLOBE_Regions_df$files$filename == 'ArchaeoGLOBE_Simplified_Regions.zip', ]

# read into local dir
walk(ArchaeoGLOBE_Regions_df_files$label,
     ~get_file(.x, ArchaeoGLOBE_Regions_DOI) %>%
       writeBin(paste0('data/raw-data/', .x)))

unzip('data/raw-data/ArchaeoGLOBE_Simplified_Regions.zip',
      overwrite = TRUE,
      exdir = 'data/raw-data/ArchaeoGLOBE_Simplified_Regions')

# read into the current environment, and simplify the polygons for faster plotting
regions <- st_read('data/raw-data/ArchaeoGLOBE_Simplified_Regions/ArchaeoGLOBE_Simplified_Regions.shp',
          quiet = TRUE)
```
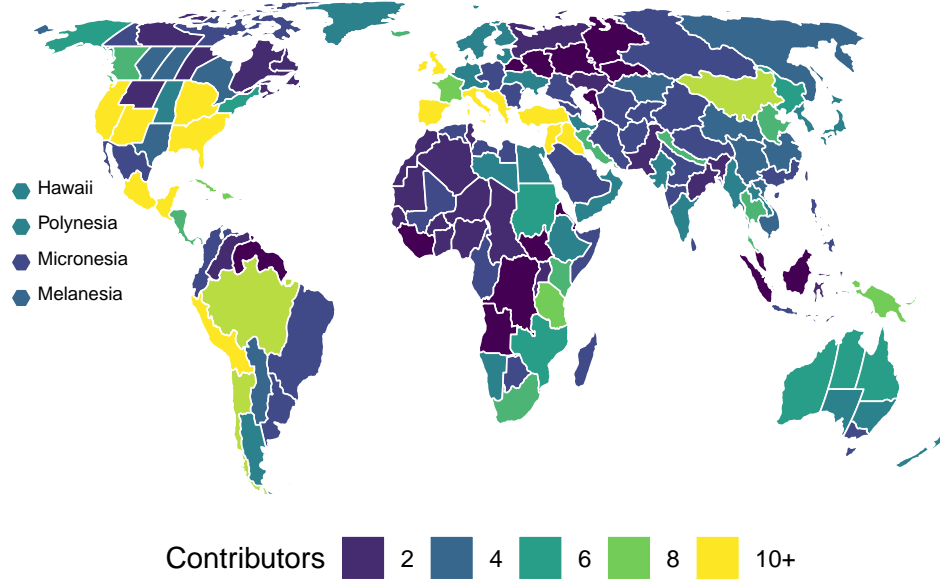
Exploratory plots

Before running any analyses, let's look at the data. How many responses do we have per region?
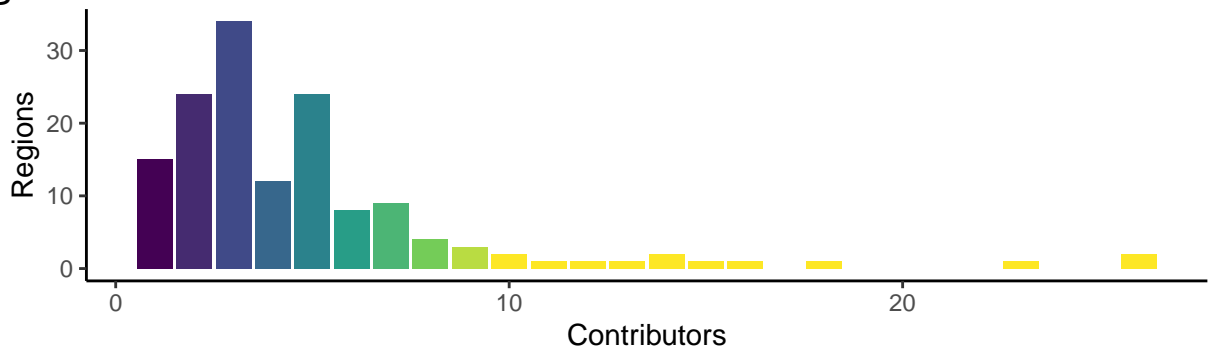
```r
response_counts <- archaeoglobe %>%
  group_by(REGION_ID) %>%
  count %>%
  mutate(n10 = replace(n, n > 10, 10))
```

A

B

How many published archaeolgical excavations are estimated for each region?

```r
# a function for calculating the mode, from https://stackoverflow.com/a/4684647
calculate_mode <- function(x) {
  uniqx <- unique(x)
  uniqx[which.max(tabulate(match(x, uniqx)))]
}

# use this vector to order the RN_SITES variable by increasing number of sites
site_order <- c('< 50', '50-249', '250-499', '500-999', '> 1000')

# find the modal response for the number of published excavations in each region
site_counts <- archaeoglobe %>%
  select(REGION_ID, RN_SITES) %>%
  group_by(REGION_ID) %>%
  summarise(sites = calculate_mode(RN_SITES)) %>%
  mutate(sites = ordered(sites, levels = site_order))
```

Published Excavations

Legend:
- Hawaii
- Polynesia
- Micronesia
- Melanesia

< 50 50–249 250–499 500–999 > 1000
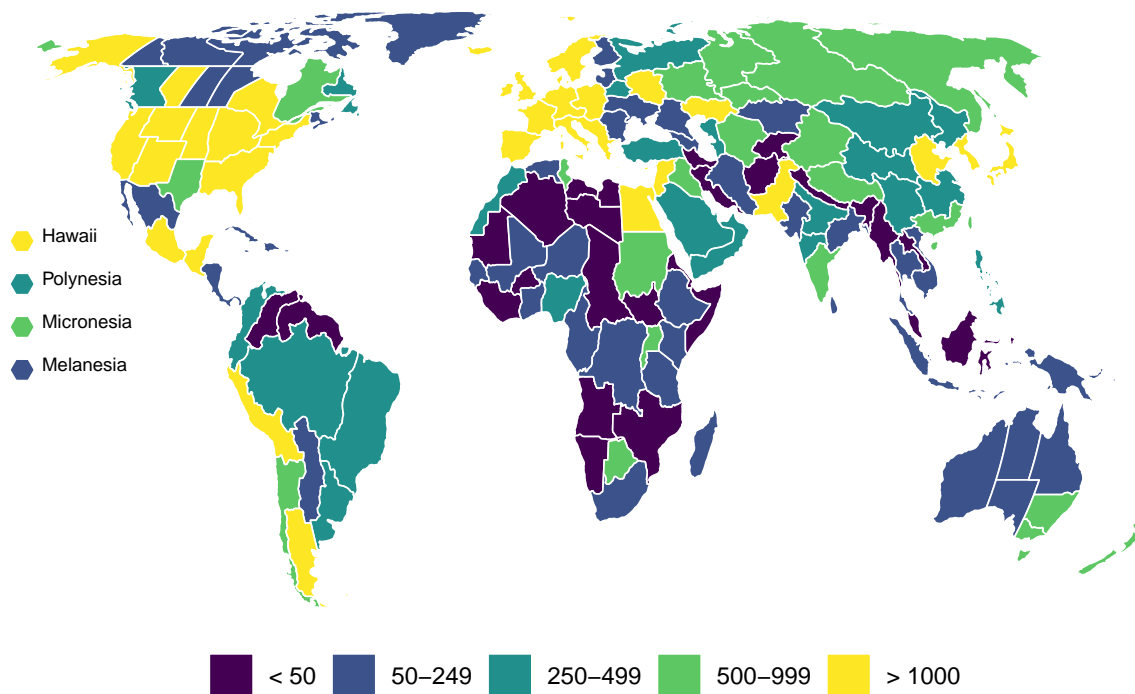
Analysis functions

Define some analysis functions that we'll be using repeatedly in the analysis, so that we don't have to keep copying and pasting the same lines of code.

This function subsets the data to highlight a variable of interest, and converts it from a wide to a long "tidy" format to make analysis and plotting easier.

```
preprocess <- function(prefix, categories){
  archaeoglobe %>% # start with the full ArcheoGlobe data
    # drop columns not related to the variable of interest
    select(c(CONTRIBUTR:LAND_AREA, starts_with(prefix))) %>%
    gather(time, value, starts_with(prefix)) %>% # one value per row
    mutate(time = parse_number(time) * -1, # convert time period labels to years
           value = ordered(value, levels = categories),
           cat_num = as.numeric(value)) %>%
    mutate_if(is.character, as.factor) # convert characters to factors
}
```

This function takes a data frame produced by the above function and fits GAM to the global trend and local deviations for each region, accounting for inter-observer variability. This function takes as arguments a preprocessed data frame containing time slices, regions, contributors, and the ordered categorical response variable transformed to a numeric vector.

```
cores <- max(parallel::detectCores() / 2, 1) # physical cores for parallelization
cl <- parallel::makeCluster(cores)
```

```
fit_gam <- function(x, n_cats){
  bam(cat_num ~
        # this spline is for the global trend
        s(time, bs = 'cr', m = 2) +
        # region-specific trends. bs = 'ts' and m = 1
        # help penalize deviation from the global model
        s(time, by = REGION_LAB, bs = 'cs', m = 1) +
        # add back in region-specific intercepts
        REGION_LAB  +
        # model contributor as a random effect
        s(CONTRIBUTR, bs = 're'),
      data = x, # data frame to analyize
      family = ocat(R = n_cats), # ordered categorical with n levels
      # final 3 arguments just speed up the model fitting
      method = 'fREML',
      discrete = TRUE,
      cluster = cl)
}
```

This function extracts the estimated trends for each region, incorporating the global and regional splines as
well as the region and contributor specific intercepts. Then it clusters these trends into 6 discrete clusters
using k-means. The choice of 6 clusters is somewhat arbitrary, and is made simply based on visual comparisons
of different cluster solutions with the goal of ensuring visually interpretable results.

```
extract_trends <-function(mod, n_clusters = 6){
  set.seed(1000) # set seed for reproducability of clusters
  archaeoglobe %>% # create dummy data for prediction in the following lines
    select(REGION_LAB) %>%
    group_by(REGION_LAB) %>%
    slice(1) %>%
    slice(rep(1:n(), each = 198)) %>%
    ungroup %>%
    mutate(time = rep_len(seq(-10000, -150, 50), n()),
           CONTRIBUTR = 'CYRBU') %>% # select an arbitrary contributor
    mutate(preds = predict(mod, .)) %>% # estimate trend lines
    mutate(preds = plogis(preds)) %>% # transform responses to [0,1] scale
    spread(time, preds) %>%
    # next is the actual kmeans clustering code
    mutate(cluster = kmeans(.[,-c(1,2)], n_clusters, iter.max = 100, nstart = 100)$cluster)
}
```

Analysis

GAMM trends

Now we use the functions defined above to estimate trends in ArchaeoGlobe data. For convenience, first
define a data frame that lists the prefixes of the variables we are interested in (e.g. "EXP" for expertise) and
the levels of the ordered factors associated with each variable. This will make it easier to quickly focus on a
specific variable. The `tribble` command is simply a way to make a data frame by row rather than column,
which makes the code easier to read.

```
response_levels <- tribble(
  ~prefix, ~categories,
  'EXP', c('None', 'Low', 'High'),                    # Expertise
  'DQ', c('Unknown', 'Low', 'Moderate', 'Good'),      # Data Quality
  'HUNT', c('none', 'minimal (<1%)', 'common (1-20%)', 'widespread (>20%)'),
  'EXAG', c('none', 'minimal (<1%)', 'common (1-20%)', 'widespread (>20%)'),
  'INAG', c('none', 'minimal (<1%)', 'common (1-20%)', 'widespread (>20%)'),
  'PAST', c('none', 'minimal (<1%)', 'common (1-20%)', 'widespread (>20%)'),
  'URBN', c('Absent', 'Present')
)
```

Now map each of the above functions to each variable. This allows us to run the analysis for all variables of interest in a single step, and save all the outputs in a tibble format for easy plotting. If you're running this for the first time, it should take about 40 minutes to run on a Intel NUC with a 5th-gen Intel Core i7-5557U processor and 16gb of RAM running Linux. By default, we pull pre-computed results from a repository rather than running the time consuming analysis.

```
trend_dat <- response_levels %>%
  mutate(data = map2(prefix, categories, ~preprocess(.x,.y)),
         n_cats = map_dbl(categories, length),
         mod = map2(data, n_cats, fit_gam),
         trends = map(mod, extract_trends))

# save to disk
saveRDS(trend_dat, file = 'data/derived-data/trend_dat.rda')

# write a note to indicate the provenance of this file
this_commit <-  git2r::revparse_single(git2r::repository('.'), "HEAD")

writeLines(paste0('trend_dat generated on ',
                  Sys.Date() , ' from archaeoglobe.Rmd at git commit ',
                  this_commit$sha, ' made by ',
                  this_commit$author$name, ' on ',
                  this_commit$author$when, " with the message '",
                  this_commit$summary, "'"),
           con = 'data/derived-data/README.md',
           sep = '')
message('Done.')
```

Results

First we plot out the global trends for each land use type, and compare them to the consensus estimates. Then we plot the local (regional trends) for all land use types, and map out their associated clusters. Please refer to the .rmd source file for the code to make the plots.

The global trend in foraging shows constant high prevalence until around 6,000 years ago, after which there is a smooth decline until the present day when it is very rare. Mapping out the clusters reveals a clear east-west divide, which regions in Afro-eurasia seeing foraging earlier then the global mean, and regions in the Americas and Oceania seeing later peaks in foraging.

The global trends in the prevalence of pastoralism, extensive and intensive agriculture, and urbanism all follow a sigmoidal curve, which means the trend is linear on the scale of the linear predictor (the ordered

6

categorical GAM uses a logit transform as a latent link function). This means that there is a simple increase in the probability of each land use type being prevalent over time.

How does self-professed level of expertise vary in each region over time? The global trend is a roughly linear increase in self-reported expertise from 10ka BP up to 2ka BP, then a falloff continuing to the present day. The present day expertise values are approximately the same as at 10ka BP. This makes sense, as it points to both the increased frequency of preserved archaeological materials with time as well as the reduction in archaeological attention in periods with extensive historical records.

The global trend in data quality is more or less the same as the expertise data, with the peak in data quality occurring more recently than for expertise and with a less dramatic falloff leading to the present day. Unlike expertise, which reaches the same values at 10ky BP and present, data quality in the present day remains high in spite of the falloff in the last 2 millennia. Also note the confidence interval for the global trend is generally wider than for the expertise responses.
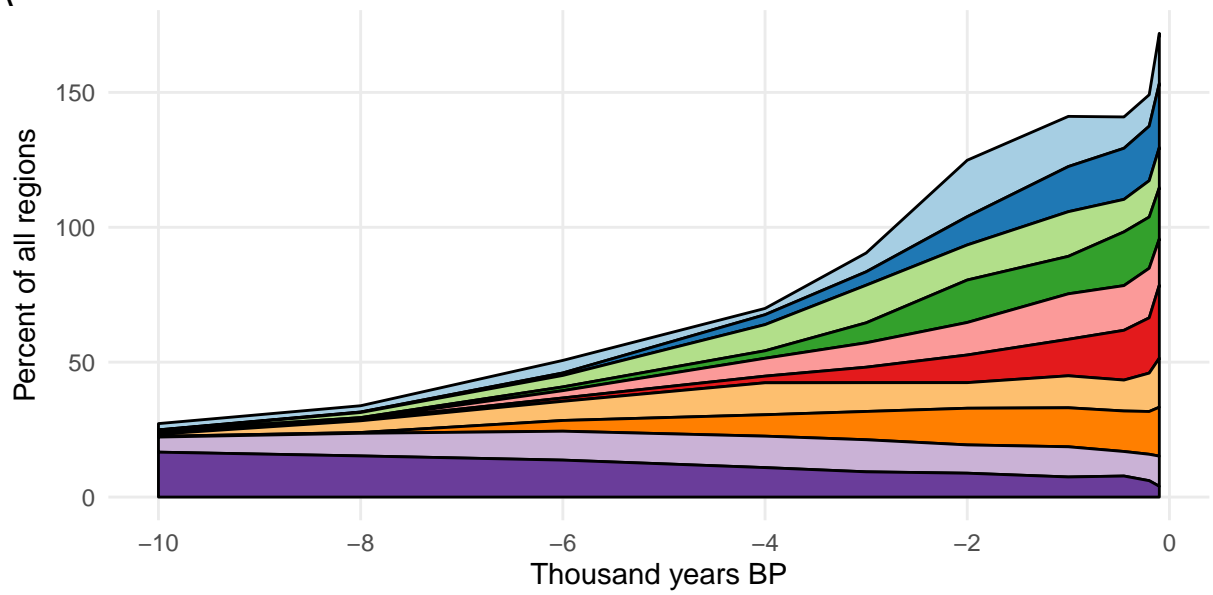
Global Trends
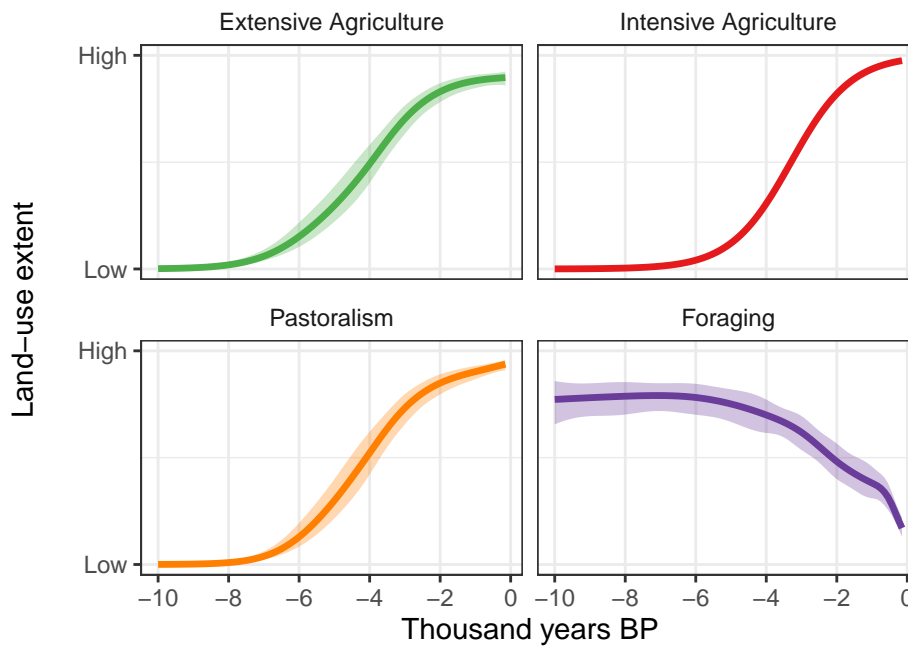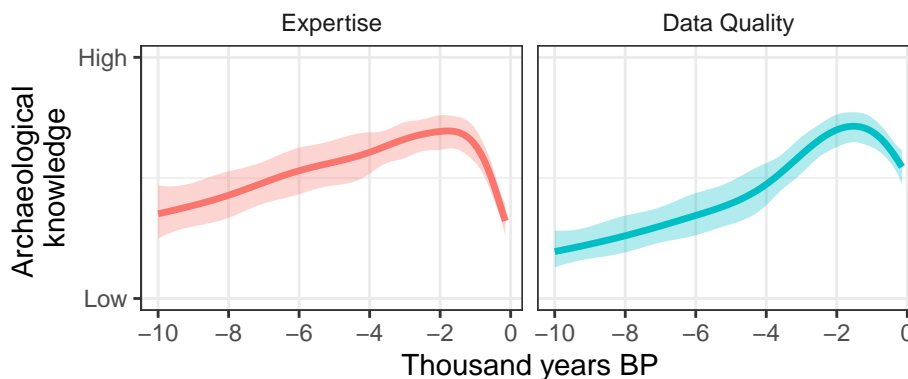
The numerical cutpoints between the ordered categorical response levels estimated by the model vary across land-use type. This is a normal result of the ordered categorical regression, and basically means that different sources of error/uncertainty impact how contributors translate their mental models of areal extent (the latent, "real" value the regression is trying to estimate) into discrete categories across the different land use types.
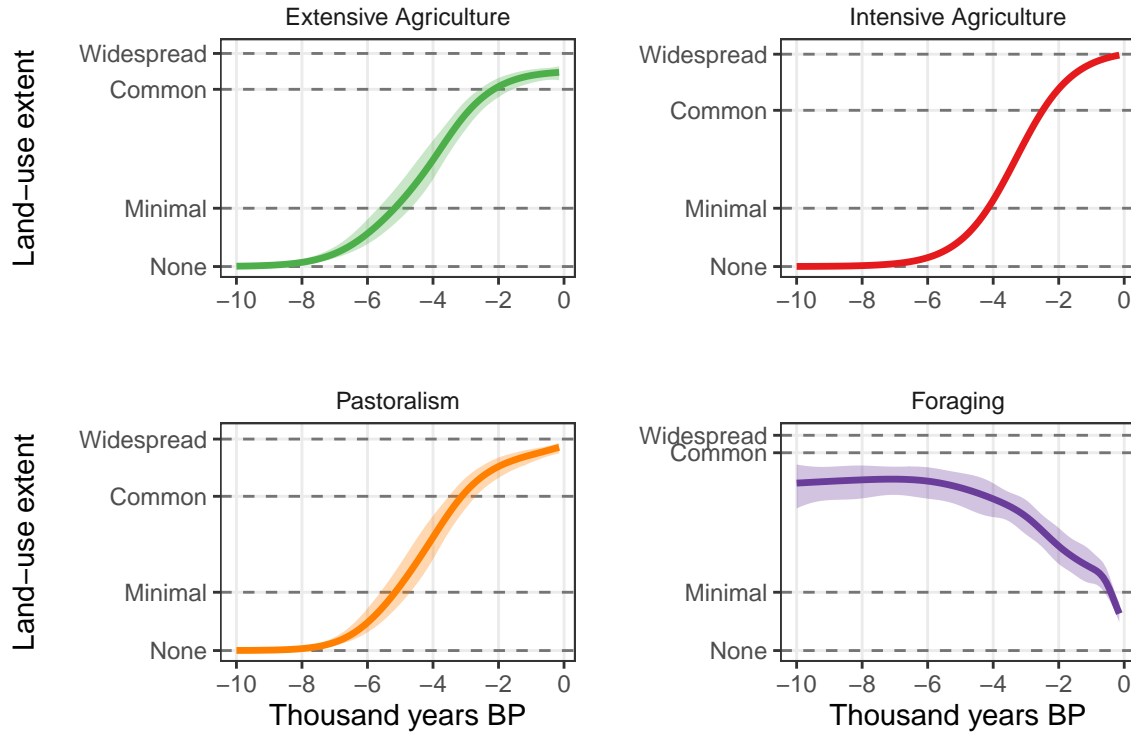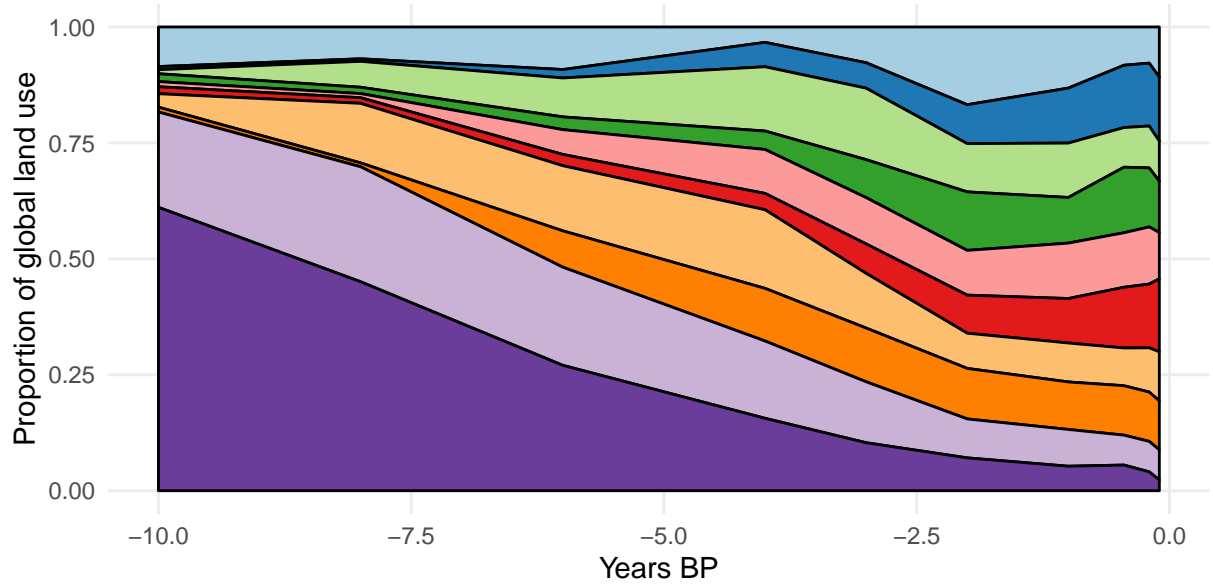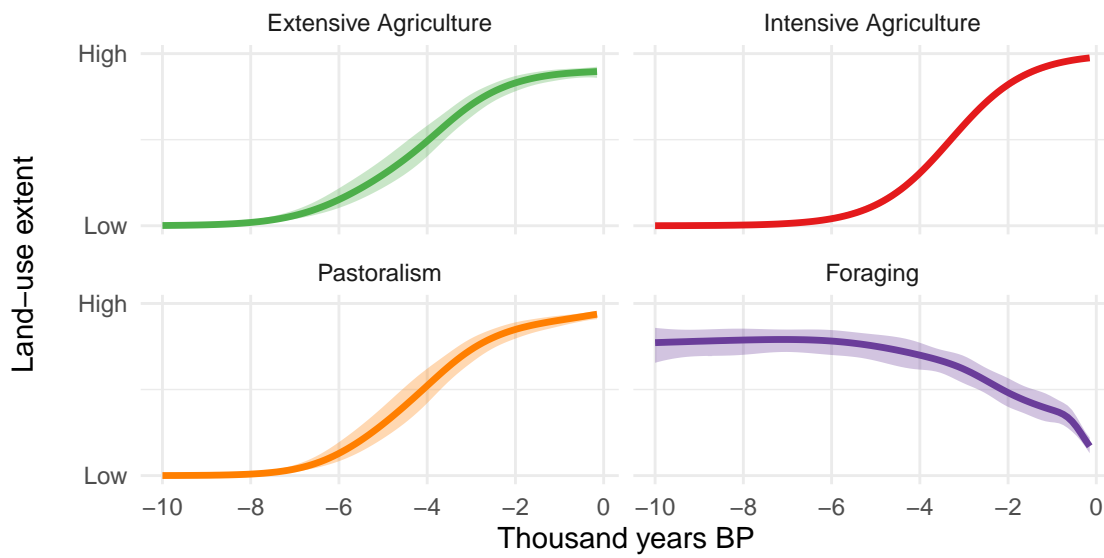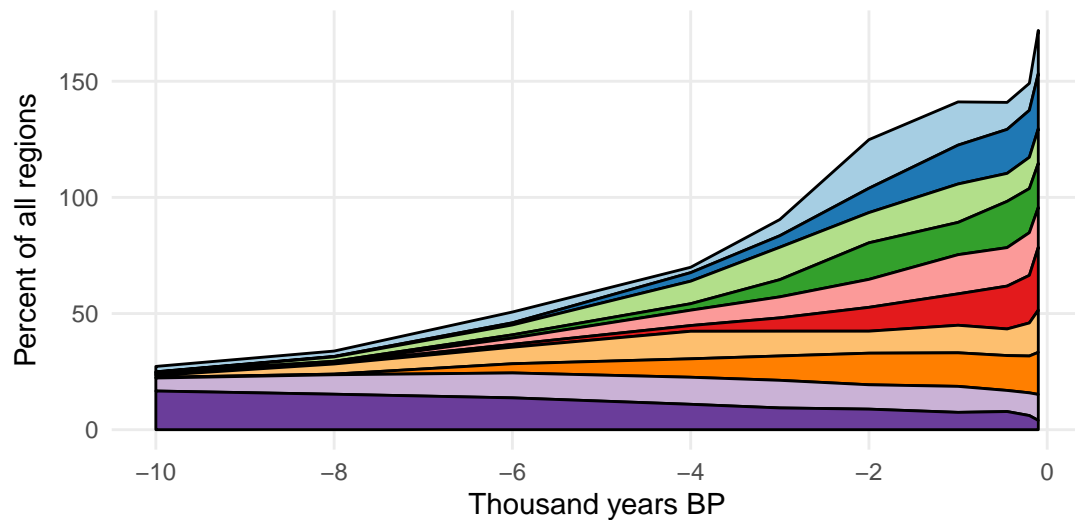

Compare the global trends to the consensus assessments.

```
(g1 + theme_minimal()) / cs1 + plot_annotation(tag_levels = 'A')
```

```
ggsave('figures/3_trends_global.png', height = 7.5, width = 6.5)
```
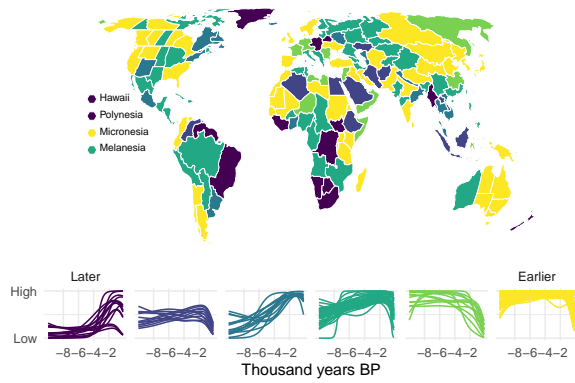
Regional Trends

A
Extensive Agriculture

B
Intensive Agriculture

C
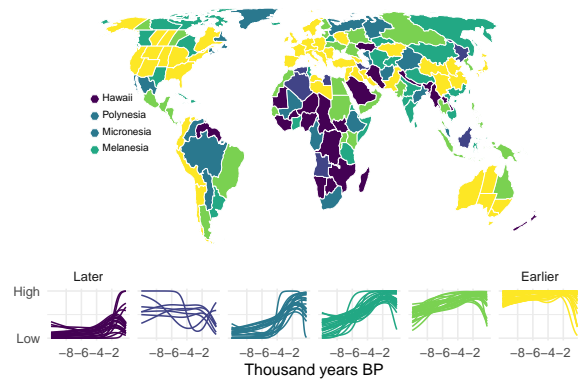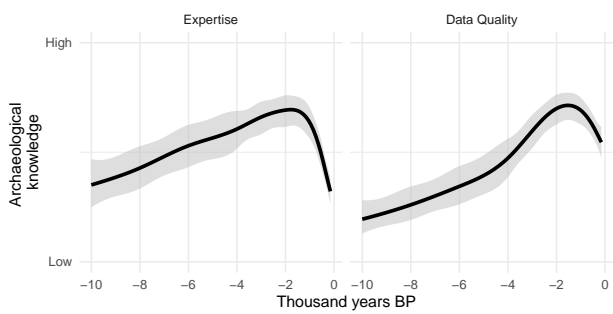### Pastoralism

D
Foraging
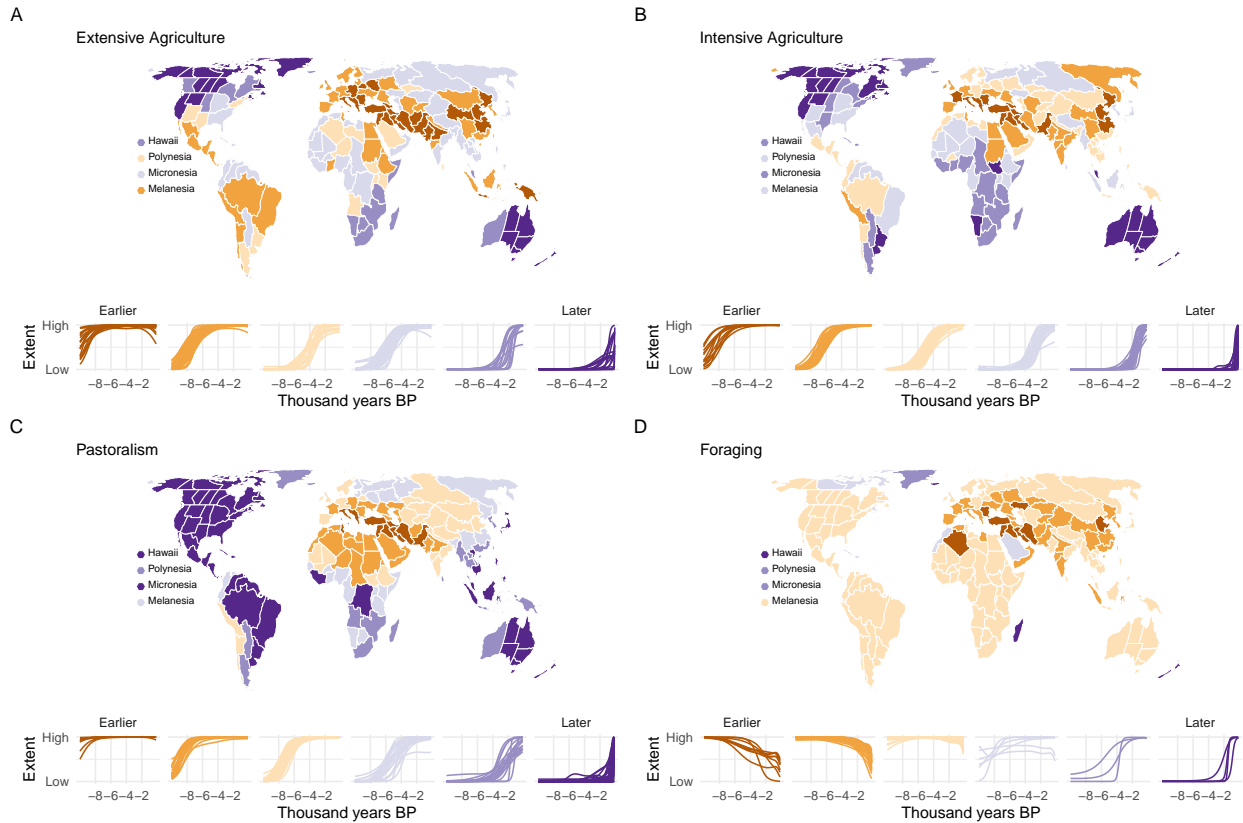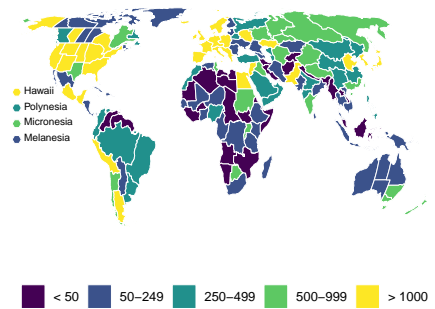
A Expertise

B Data Quality

C

D Published Excavations

```
#temporary stopping point
knitr::knit_exit()
```